

**EXCESS****Data-structures-framework**

EXCESS Concurrent Data Structures (Adapter) Library

Anders Gidenstam



Status

**EXCESS**

- Integration with the EXCESS (Search) Tree Library from UiT.
Done.
- Integrates also with NOBLE, Intel TBB (and some research prototypes)
- Available in
 - The data-structures-library repository at the EXCESS project page at GitHub:
<https://github.com/excess-project/data-structures-library>
 - src/include/EXCESS in the data-structures-framework repository at the EXCESS GitLab.

C++ library of concurrent data structures



■ Motivation

- Yet another data structures library?
 - There are many already – but if multiple choices of implementations is desired any one single library will not do.

■ Requirements

- Uniform interfaces for implementations of each ADT
- Easy to use
- Flexible
 - Should be possible to adapt existing implementations
 - Should provide what applications need
- Efficient

3

C++ library of concurrent data structures



■ Easy to use

- One shared variable for a queue instance

```
#include <EXCESS/concurrent_queue>

// Prepare a queue storing int* pointers.

// Class-wide queue pointer - can be any queue implementation.
excess::concurrent_queue<int> *any_queue_ptr;

// Pointer to specific implementation.
excess::concurrent_queue_MSTLB<int> *two_lock_queue_ptr;

// Or one instance of a specific implementation.
excess::concurrent_queue_MSTLB<int> two_lock_queue;
```

- Run-time selection of implementation

```
void init(int queue_to_use)
{
    switch (queue_to_use) {
        case 7:
            // The concurrent_queue from Intel Threading Building Blocks.
            any_queue_ptr = new excess::concurrent_queue_TBBQueue<int>();
            break;
        default:
            // An implementation of the Michael and Scott two-lock queue.
            any_queue_ptr = new excess::concurrent_queue_MSTLB<int>();
    }
}
```

4

C++ library of concurrent data structures



■ Easy to use

- To use (perform operations) on it a handle is needed
 - Some implementations need to know which/how many threads are going to use it
 - Making this handle explicit is more efficient than trying to hide it

```
void foo<concurrent_queue_t>(concurrent_queue_t* queue_to_use)
{
    #pragma omp parallel
    {
        // Register this thread with the concurrent queue (e.g. for memory
        // management). It is more efficient to do this explicitly rather than
        // checking some thread local variable internally on every call.
        concurrent_queue_t::handle* queue_handle = queue_to_use->get_handle();

        for (int i = 0; i < 100; ++i) {
            queue_handle->enqueue(new int(i));
        }

        int* tmp;
        while (queue_handle->try_dequeue(tmp)) {
            cout << "Got '" << *tmp << "'." << std::endl;
            delete tmp;
        } // Exit when the queue is empty.
        // Deregister this thread from the queue.
        delete queue_handle;
    } // end of #pragma omp parallel
}
```

5

Concurrent Producer/Consumer Collections



■ The **concurrent_queue<T>** interface

- Linearizeable operations
 - void enqueue(T* item)
 - bool try_dequeue(T*& item)
 - bool empty()

■ The **concurrent_stack<T>** interface

- Linearizeable operations
 - void push(T* item)
 - bool try_pop(T*& item)
 - bool empty()

■ The **concurrent_bag<T>** interface

- Linearizeable operations
 - void insert(T* item)
 - bool try_remove_any(T*& item)
 - bool empty()

6

Concurrent Producer/Consumer Collections



■ Implementations

■ Queues

■ Internal

- Two-lock queue [Michael & Scott, 1996]
- STL vector + OpenMP lock

■ NOBLE (external dependency)

- L-F queue DB [Michael & Scott, 1996]
- L-F queue DU [Valois, 1994]
- L-F queue SB [Tsigas & Zhang, 2001]
- L-F queue Basket [Hoffman, Shalev & Shavit, 2007]
- L-F queue Elim [Michael & Scott, 1996] + Elimination [Moir, Nussbaum, Shalev & Shavit, 2005]
- L-F queue BB [Gidenstam, Sundell, Tsigas, 2010]
- L-B queue

■ Intel TBB (external dependency)

- Concurrent_queue

7

Concurrent Producer/Consumer Collections



■ Implementations

■ Stacks

■ NOBLE (external dependency)

- L-F stack B [Michael, 2004]
- L-F stack Elim [Michael, 2004] + Elimination [Hendler, Shavit & Yerushalami, 2010]

■ (Unordered) Bags / Pools

■ NOBLE (external dependency)

- L-F bag [Gidenstam, Sundell, Papatriantafilou & Tsigas, 2011]
- L-F pool EDTree [Afek, Korland, Natanzon & Shavit, 2010]

■ (+ all queues and stacks are also admissible as bags.)

8

Concurrent Dictionaries



- The `concurrent_weak_dictionary<Key, T>` interface
 - “Weak” means the semantics of the operations impose few(er) consistency demands.
 - Linearizeable operations
 - `void insert(Key key, T* value)`
 - Insert the key-value pair.
 - NOTE: Returns no information on whether the key existed before or not.
 - `bool lookup(Key key, T*& value)`
 - Sets value to the associated value and returns true if key exists in the dictionary. Returns false otherwise.
 - `void remove(Key key)`
 - Removes the key-value association from the dictionary.
 - NOTE: Returns no information on whether the key existed before or not.

9

Concurrent Dictionaries



- Implementations
 - Search tree algorithms
 - The EXCESS search tree library (internal dependency)
 - GreenBST
 - DeltaTree
 - CBTree
 - Hash table algorithms
 - Intel TBB (external dependency)
 - hashmap
 - Various research prototypes (often less than stable) (external dependency)
 - L-F Cuckoo hashing, [Nhan & Tsigas, 2014]
 - L-F bucketized Cuckoo hashing, [Nhan & Tsigas, 2014]
 - Hopschotch hashing, [Herlihy, Shavit & Tzafrir, 2008]
 - Bitmapped Hopschotch hashing, [Herlihy, Shavit & Tzafrir, 2008]
 - L-B chained, [Lea, ?] (alg. From `java.util.concurrent`, implemented in the hopschotch microbenchmark)

10

The Benchmark framework



■ Front-end testbench program

- Sets and handles
 - #threads
 - Pinning strategy
 - Duration
 - Parsing per-experiment command line parameters
 - Formatting output
- Integration with
 - The ATOM monitoring framework
 - MeterPU for power (neither committed nor actually tested yet – Intel PCM needs root access and doesn't work on the EXCESS server at Chalmers)
- Experiments (hierarchy of classes added at compile time)
 - Producer-Consumer microbenchmark (uses concurrent_bag)
 - Mandelbrot application (uses concurrent_bag)
 - SGEMM microbenchmark (uses concurrent_bag)
 - SpDGEMM microbenchmark (uses concurrent_bag)
 - Dictionary microbenchmark
 - Weak dictionary microbenchmark (uses concurrent_weak_dictionary)